

RNN from scratch

Muhammad Tamjid Rahman

Contents

Loading R packages	1
1 Transformers and Attention	1
2 Implementing a simple RNN	3
2 Implementing a simple RNN	6

Loading R packages

```
library(uuml)
```

1 Transformers and Attention

Loading dataset

```
data("transformer_example")
```

1)

```
# Picking out the matrix for the first attention head
```

```
Wq <- transformer_example$Wq[, , 1]
```

```
Wk <- transformer_example$Wk[, , 1]
```

```
Wv <- transformer_example$Wv[, , 1]
```

```
# Picking out the first three words and their embeddings
```

```
X <- transformer_example$embeddings[1:3,]
```

Implementing the function qkv()

```
qkv = function(X, Wq, Wk, Wv){  
  Q=X%*%Wq  
  K=X%*%Wk  
  V=X%*%Wv  
  return(list(Q=Q,K=K,V=V))  
}
```

```
qvk(X,Wq,Wk,Wv)
```

```
## $Q
##           [,1]      [,2]      [,3]
## the      0.4722259  0.04995783 -0.5350845
## quick -0.3662435  0.12144160  0.3454785
## brown -0.1029677 -0.12728414  0.1817097
##
## $K
##           [,1]      [,2]      [,3]
## the      0.094360579 -0.203807092 -0.1851229
## quick -0.033313240  0.279012100  0.2530560
## brown -0.004457052  0.001013468  0.0133802
##
## $V
##           [,1]      [,2]      [,3]
## the      0.317318525 -0.35023010  0.13284078
## quick  0.009929565  0.04208206 -0.15412097
## brown -0.316413241  0.27717408  0.02725089
```

2 Computing the attention

```
res=qvk(X,Wq,Wk,Wv)
```

```
# Softmax activation function
```

```
softmax <- function(qk){
  val <-exp(qk)/sum(exp(qk))

  return(val)
}
```

```
attention = function(Q,K,V){
  attention = matrix(0,nrow(K),nrow(K))
  for (i in 1:nrow(K)) {
    attention[i,] = softmax((Q%*%t(K))[i,]/sqrt(nrow(K)))
  }
  Z = attention%*%V
  return(list(Z = Z, attention = attention))
}
```

```
attention(res$Q, res$K, res$V)
```

```
## $Z
##           [,1]      [,2]      [,3]
## [1,]  0.012395453 -0.0212420459  0.009404870
## [2,] -0.003759269 -0.0008360029 -0.005108890
## [3,]  0.002412222 -0.0088974612  0.001147999
##
## $attention
##           [,1]      [,2]      [,3]
## [1,]  0.3601932  0.3080896  0.3317172
## [2,]  0.3088780  0.3582373  0.3328847
## [3,]  0.3300375  0.3360583  0.3339042
```

3 Interpreting

Attention values describes how much the attention of a word to the other word in a sequence. In the second row, the work 'quick' has 31% attention to the word 'the', 36% at its own and 33% to 'brown'.

4 Implementing a multi-head attention layer

```
multi_head_self_attention = function(X,Wq,Wk,Wv,W0){
  Z = matrix(0,ncol(Wk),1)

  for (i in 1:dim(Wk)[3]) {
    Wq1 <- Wq[, ,i]
    Wk1 <- Wk[, ,i]
    Wv1 <- Wv[, ,i]

    Q = qvk(X,Wq1,Wk1,Wv1)$Q
    V = qvk(X,Wq1,Wk1,Wv1)$V
    K = qvk(X,Wq1,Wk1,Wv1)$K

    Z = cbind(Z,attention(Q,K,V)$Z)
  }
  return(Z[,-1]%*%W0)
}

multi_head_self_attention(X,
                          transformer_example$Wq,
                          transformer_example$Wk,
                          transformer_example$Wv,
                          transformer_example$W0)
```

```
##           [,1]           [,2]           [,3]
## [1,] -0.014189613 -0.0040299008 -0.006756286
## [2,] -0.009963516 -0.0010724342 -0.001996524
## [3,] -0.006394562 -0.0006626115 -0.002219108
```

2 Implementing a simple RNN

Loading Data

```
data("rnn_example")
```

1

```
hidden_dim=4
output_dim=3

rnn_unit=function(h_t_minus_one,X,W,U,b){
  a_t=b+W*%*%h_t_minus_one+U*%*%t(X)
  h_t_minus_one=tanh(a_t)
  return(a_t=a_t)
}
```

```

X <- rnn_example$embeddings[1,,drop=FALSE]
h_t_minus_one <- matrix(0, nrow = hidden_dim, ncol = 1)
a_t <- rnn_unit(h_t_minus_one, X,
               W = rnn_example$W,
               U = rnn_example$U,
               b = rnn_example$b)

```

```
a_t
```

```

##           the
## [1,]  0.5819145
## [2,] -2.2686535
## [3,] -0.6410312
## [4,]  1.4891931

```

2 Implementing the tanh() activation function

```

activation=function(a_t){
  h_t=tanh(a_t)
  return(h_t=h_t)
}

```

```

h_t <- activation(a_t)
h_t

```

```

##           the
## [1,]  0.5240555
## [2,] -0.9788223
## [3,] -0.5656013
## [4,]  0.9031762

```

3 Implementing the output function and the softmax function

```

output_rnn=function(h_t,V,c){
  o_t= c+ V%*%h_t
  return(output_rnn=o_t)
}

```

```

yhat_t <- softmax(output_rnn(h_t, rnn_example$V, rnn_example$c))
yhat_t

```

```

##           the
## [1,] 0.3063613
## [2,] 0.2930885
## [3,] 0.4005502

```

4 Implementing the full recurrent layer

```

rnn_layer=function(X,W,V,U,b,c){

  h_t_minus_one <- matrix(0, nrow = hidden_dim, ncol = 1)
  h_t=matrix(0, nrow = hidden_dim,1)
  yhat=matrix(0,nrow = output_dim,1)

```

```

for (i in 1:nrow(X)) {
  a_t <- rnn_unit(h_t_minus_one, t(X[i,]),W,U,b)

  o_t=output_rnn(activation(a_t), V, c)

  h_t=cbind(h_t,activation(a_t))

  yhat=cbind(yhat, softmax(o_t))

  h_t_minus_one=tanh(a_t)
}

return(list(h_t=t(h_t[, -1]), yhat=t(yhat[, -1])))
}

X <- rnn_example$embeddings[1:3,,drop=FALSE]
rnn_layer(X,
  W = rnn_example$W,
  V = rnn_example$V,
  U = rnn_example$U,
  rnn_example$b,
  rnn_example$c)

## $h_t
##           [,1]           [,2]           [,3]           [,4]
## [1,]  0.52405551 -0.97882227 -0.5656013  0.9031762
## [2,] -0.05951368  0.03988226  0.8241800 -0.6562744
## [3,] -0.08984008  0.92822217 -0.1563247 -0.6657626
##
## $yhat
##           [,1]           [,2]           [,3]
## [1,]  0.3063613  0.2930885  0.4005502
## [2,]  0.2838013  0.3490452  0.3671536
## [3,]  0.2878002  0.3666877  0.3455121

```

5 The hidden state h_t for the token dog

```

X <- rnn_example$embeddings[drop=FALSE]

h_t_dog=rnn_layer(X,
  W = rnn_example$W,
  V = rnn_example$V,
  U = rnn_example$U,
  rnn_example$b,
  rnn_example$c)$h_t[9,]

h_t_dog

## [1] -0.2928465  0.1813845 -0.2190118  0.2592397

```

2 Implementing a simple RNN

Loading Data

```
data("rnn_example")
```

1

```
hidden_dim=4
output_dim=3

rnn_unit=function(h_t_minus_one,X,W,U,b){
  a_t=b+W%*%h_t_minus_one+U%*%t(X)
  h_t_minus_one=tanh(a_t)
  return(a_t=a_t)
}

X <- rnn_example$embeddings[1,,drop=FALSE]
h_t_minus_one <- matrix(0, nrow = hidden_dim, ncol = 1)
a_t <- rnn_unit(h_t_minus_one, X,
               W = rnn_example$W,
               U = rnn_example$U,
               b = rnn_example$b)

a_t

##           the
## [1,]  0.5819145
## [2,] -2.2686535
## [3,] -0.6410312
## [4,]  1.4891931
```

2 Implementing the tanh() activation function

```
activation=function(a_t){
  h_t=tanh(a_t)
  return(h_t=h_t)
}

h_t <- activation(a_t)
h_t

##           the
## [1,]  0.5240555
## [2,] -0.9788223
## [3,] -0.5656013
## [4,]  0.9031762
```

3 Implementing the output function and the softmax function

```
output_rnn=function(h_t,V,c){
  o_t= c+ V%*%h_t
```

```

    return(output_rnn=o_t)
}

yhat_t <- softmax(output_rnn(h_t, rnn_example$V, rnn_example$c))
yhat_t

##           the
## [1,] 0.3063613
## [2,] 0.2930885
## [3,] 0.4005502

```

4 Implementing the full recurrent layer

```

rnn_layer=function(X,W,V,U,b,c){

  h_t_minus_one <- matrix(0, nrow = hidden_dim, ncol = 1)
  h_t=matrix(0, nrow = hidden_dim,1)
  yhat=matrix(0,nrow = output_dim,1)

  for (i in 1:nrow(X)) {
    a_t <- rnn_unit(h_t_minus_one, t(X[i,]),W,U,b)

    o_t=output_rnn(activation(a_t), V, c)

    h_t=cbind(h_t,activation(a_t))

    yhat=cbind(yhat, softmax(o_t))

    h_t_minus_one=tanh(a_t)
  }

  return(list(h_t=t(h_t[,-1]),yhat=t(yhat[,-1])))
}

X <- rnn_example$embeddings[1:3,,drop=FALSE]
rnn_layer(X,
          W = rnn_example$W,
          V = rnn_example$V,
          U = rnn_example$U,
          rnn_example$b,
          rnn_example$c)

## $h_t
##           [,1]      [,2]      [,3]      [,4]
## [1,] 0.52405551 -0.97882227 -0.5656013 0.9031762
## [2,] -0.05951368 0.03988226 0.8241800 -0.6562744
## [3,] -0.08984008 0.92822217 -0.1563247 -0.6657626
##
## $yhat
##           [,1]      [,2]      [,3]
## [1,] 0.3063613 0.2930885 0.4005502
## [2,] 0.2838013 0.3490452 0.3671536
## [3,] 0.2878002 0.3666877 0.3455121

```

5 The hidden state h_t for the token dog

```
X <- rnn_example$embeddings[drop=FALSE]
```

```
h_t_dog=rnn_layer(X,  
  W = rnn_example$W,  
  V = rnn_example$V,  
  U = rnn_example$U,  
  rnn_example$b,  
  rnn_example$c)$h_t[9,]
```

```
h_t_dog
```

```
## [1] -0.2928465  0.1813845 -0.2190118  0.2592397
```

The role of the parameters U , W in the recurrent net

U is the weight matrix for the input vector x and W is the weight matrix for the state of previous time step in hidden layer (eq 10.8).¹ In our model in `simple_rnn` (SimpleRNN) layer the dimension of W is 32x16 and for U is 16x16.

The role of the parameters V

The parameter V is the weight matrix for the hidden-to-output connections (eq. 10.10).¹ In our model in `dense_2` (Dense) layer it is included. The dimension of V is 32x32.